

SEPTEMBER 2026 / EDITION 1.0


AI in Indian Finance

Where it is actually creating value

THE QUESTION

How much authority should the model have?

The higher the consequence of a wrong decision, the stronger the case for provenance, human review and clear accountability.



VALUE

5 workflow lanes

READ 03

EVIDENCE

Regulators + annual reports

READ 09

CONTROL

Model risk + oversight

READ 10

Useful AI is mostly workflow AI

Ignore the science-fiction framing. In Indian finance, the practical story is already visible in the operating layer.

Five places to look

- Customer service: faster routing, retrieval and first-line resolution.
- Fraud and surveillance: pattern detection that narrows what humans need to inspect.
- Lending and credit: document processing and decision support around underwriting.
- Research and compliance: summarisation, monitoring, document review and analyst assistance.
- Operations: reconciliation, claims, voice-to-data conversion and employee copilots.

The control question

Automation can reduce friction without removing accountability. This issue separates low-consequence assistance from decisions that can deny credit, freeze money, shape advice or materially affect a customer. [\[S4\]](#) [\[S6\]](#) [\[S8\]](#)

Evidence note

Regulatory documents are treated as primary policy evidence. Company annual reports are used only for clearly labelled company-reported operating metrics. Drafts and consultations are labelled as drafts or consultations, not final rules.

03	The value map
04	Customer service
05	Fraud and surveillance
06	Lending and credit
07	Research, advice and compliance
08	Operations
09	India's governance direction
10	Model risk and human oversight
11	Practical control checklist
12	Sources and methodology


Read the issue as a map of value plus control: what AI can accelerate, what must remain reviewable, and who remains responsible.

Five workflows, one pattern

The recurring pattern is not autonomous finance. It is AI reducing the amount of information a person or system must manually process before the next decision.

MAP 01 / WHERE PRIMARY SOURCES SHOW AI AT WORK

NOT A MARKET-SIZE ESTIMATE



The useful middle

SEBI's own inventory of market use cases includes surveillance, KYC/AML, risk management, data analysis, document review, cyber security, customer service and chatbots. RBI examples include fraud detection and complaint handling. [\[S2\]](#) [\[S5\]](#) [\[S7\]](#)

Where value compounds

AI becomes more useful when the input is high-volume, repetitive and reviewable: documents, conversations,

transaction patterns, complaints, filings and internal knowledge. That does not make the downstream decision low-risk.

What this chart is - and is not

The visual is a MarketDeck workflow map, not a ranking of returns or an estimate of market size. It shows where primary sources document actual use or deployment.


Editorial takeaway: the strongest evidence today is around throughput and decision support, not autonomous financial judgment.

The first win is response capacity

Service is a natural entry point because a bot can retrieve, classify and resolve routine work while preserving an escalation path for higher-stakes cases.

CHART 01 / COMPANY-REPORTED SERVICE METRICS

FY2026 · DIFFERENT BUSINESSES / DENOMINATORS



RBI complaint management
Phase I: conversational AI chatbot for complainants. Phase II: more advanced complaint processing.

A concrete operating metric

Bajaj Finserv reports that 65% of customer servicing at Bajaj Finance was handled through AI voice and text bots in FY2026. It also reports that Bajaj Finserv Health resolved about 55% of customer queries using agentic AI. These are company-reported operating metrics, not independent quality scores. [\[S9\]](#)

The regulator is experimenting too

RBI's 2024-25 Annual Report says it is working with ReBIT to integrate AI into the Complaint Management

System in phases: phase I for a conversational AI chatbot for complainants, followed by more advanced complaint-processing features. [\[S2\]](#)

Design the handoff, not just the bot

A service system should know when it is answering a routine question, when it is interpreting a rule, and when the customer needs a human. The practical KPI is not simply containment rate; it is resolution quality with a reliable escalation path.


Good service AI reduces waiting and repetitive work. It should not make a difficult complaint harder to reach a human with.

03 / FRAUD & SURVEILLANCE

AI is useful when it narrows the haystack

Fraud systems are a strong fit for machine learning because the first job is often to surface unusual behaviour quickly - not to declare guilt.

SCHEMATIC 01 / DETECTION IS NOT ADJUDICATION



FALSE POSITIVE COSTS

Blocked payments · delayed access · extra verification · customer distress

MuleHunter.AI is a useful example

RBI described MuleHunter.AI, developed by RBI Innovation Hub, as an AI/ML model for detecting mule bank accounts more efficiently. A pilot with two large public sector banks produced encouraging results, and banks were encouraged to collaborate further. [\[S3\]](#)

Surveillance is already a securities-market use case

SEBI's board material lists surveillance, KYC/AML, risk management, behaviour analysis and cyber security among

observed AI/ML uses across brokers, depositories, market infrastructure institutions and mutual funds. [\[S5\]](#)

The decision boundary matters


A model can rank suspicious accounts or transactions without becoming the final authority. False positives have real costs: blocked payments, delayed access, additional verification and customer distress. The higher the impact, the stronger the case for documented review and redress.

Detection is not adjudication. The safest architecture separates anomaly scoring from the decision that affects the customer.


Faster underwriting is valuable - but consequential

Lending combines exactly what AI likes - documents, repeated decisions and large datasets - with exactly what governance worries about: fairness, explainability and access to credit.

CHART 02 / TWO OPERATING SIGNALS FROM BAJAJ FINANCE COMPANY-REPORTED · FY2026



DOCUMENT PROCESSING CREDIT JUDGMENT



AI is already connected to loan workflows

Bajaj Finserv reports Rs 5,520 crore of conversational-AI-enabled loan disburseals at Bajaj Finance in FY2026 and more than 69 million application documents processed using Vision AI. 'AI-enabled' should not be read as 'AI made the credit decision by itself'. [\[S9\]](#)

Document automation is not the same as credit judgment

Extracting data from applications, checking completeness and routing cases can remove manual work. Credit

approval, pricing and adverse decisions introduce different risks because model errors can directly change a customer's access to finance.

Model risk is moving to the centre

RBI's June 2026 draft model-risk guidance expands the frame beyond credit models to all models used by covered regulated entities, including third-party and AI/ML models, across the model lifecycle. The document is a draft released for consultation, not final regulation. [\[S4\]](#)

The more a model changes a customer's economic outcome, the less adequate 'the model said so' becomes as an explanation.

AI can compress research; responsibility does not compress

Financial research is a high-leverage use case because models can search, summarise and compare. But the output can travel into advice, execution and compliance decisions.

MATRIX 01 / DOCUMENTED SECURITIES-MARKET USE CASES

SEBI PRIMARY SOURCES

EXCHANGES

surveillance · cyber security · member support · data automation

BROKERS

KYC/docs · recommendations · chatbots · AML · order execution

MUTUAL FUNDS

customer support · surveillance · cyber security · segmentation

Regulated activity → responsibility for data, output and compliance remains with the regulated person.

SEBI documents broad use across the market

SEBI's 2025 consultation paper records AI/ML use by exchanges for surveillance, cyber security, member support and data automation; by brokers for KYC/document processing, recommendations, chatbots, account opening, surveillance, AML and order execution; and by mutual funds for customer support, surveillance, cyber security and segmentation. [\[S7\]](#)

Accountability follows the regulated activity

SEBI's 2025 Intermediaries Amendment Regulations make covered regulated persons solely responsible for investor

and stakeholder data, the outputs from AI/ML tools they rely on, and compliance with applicable law - whether the tools are built internally or sourced from third parties. [\[S6\]](#)

Advice adds disclosure duties


SEBI's January 2025 investment-adviser guidelines state that advisers using AI tools remain responsible for client data, other information used to arrive at advice, AI-based advice and legal compliance, and must disclose the extent of AI use to clients. [\[S8\]](#)

AI may shorten the path from document to insight. It does not outsource fiduciary or regulatory responsibility.

The quiet gains are in the back office

Operations is where AI can create large cumulative value without needing to pretend it can forecast markets.

CHART 03 / OPERATING SCALE, NOT INVESTMENT PERFORMANCE



~67% EXPECTED motor-claims TAT reduction through automated invoice processing - an expectation, not an achieved result.

Voice and documents become structured data

Bajaj Finserv reports 52 million voice-to-data conversions for customer conversations at Bajaj Finance in FY2026, alongside more than 69 million application documents processed using Vision AI. [\[S9\]](#)

Claims and servicing become more automatable

The same annual report says Bajaj General expects about a 67% reduction in turnaround time for motor claims through automated invoice processing. Because the figure is

explicitly an expected reduction, this issue does not present it as an achieved outcome. [\[S9\]](#)

Internal copilots can scale institutional knowledge

Axis Bank's FY2024-25 annual report says Axis Deep Intelligence, its GenAI internal chatbot, was deployed across more than 5,500 branches and supported more than 100,000 employees. That is an internal productivity deployment, not a claim about investment returns or customer outcomes. [\[S10\]](#)


The back office is a powerful proving ground: repetitive enough to automate, structured enough to measure, and close enough to human review to recover from errors.

Enable innovation, keep responsibility visible

India's current policy direction is not a simple ban-or-permit choice. It increasingly combines encouragement of AI adoption with explicit responsibility for outcomes.

FRAMEWORK 01 / FREE-AI: ENABLE + GOVERN

7 SUTRAS · 6 PILLARS · 26 RECOMMENDATIONS



FREE-AI is a framework, not a marketing slogan

RBI's FREE-AI committee report sets seven foundational Sutras, 26 actionable recommendations and six strategic pillars. The RBI says the aim is to encourage innovation in harmony with risk mitigation. [\[S1\]](#)

SEBI puts responsibility on the regulated person

Under SEBI's 2025 Intermediaries Amendment Regulations, covered intermediaries remain responsible for

data privacy/security/integrity, AI outputs they use, and legal compliance even when technology is procured from a third party. [\[S6\]](#)

Policy status matters

The FREE-AI report is a committee framework. SEBI's responsible-AI paper is a consultation paper. RBI's 2026 model-risk guidance is a draft. This issue labels each accordingly and does not describe proposals as binding final rules. [\[S1\]](#) [\[S4\]](#) [\[S7\]](#)


The emerging principle is durable: innovation can be delegated to tools; accountability cannot be delegated away.

08 / MODEL RISK & HUMAN OVERSIGHT

The hard problem begins when the model can act

A summariser and an autonomous decision engine should not share the same control model. Consequence should determine governance.

LADDER 01 / MATCH CONTROL TO DECISION AUTHORITY



MORE AUTHORITY / LOWER REVERSIBILITY →

RBI is broadening the model-risk lens

RBI's June 2026 draft guidance covers governance, risk tiering, lifecycle management, validation, continuous oversight, change management and business continuity across models used by covered regulated entities, including third-party and AI/ML models. [\[S4\]](#)

Think in levels of authority

- Assist: retrieve, classify, summarise or draft. A human remains the decision-maker.
- Recommend: score, rank or propose an action. Human review is explicit for material outcomes.

- Act within bounds: execute only inside predetermined limits with logs, overrides and monitoring.
- High-impact autonomous action: demands the strongest evidence, controls, accountability and redress.

Explainability is a control surface

The right explanation depends on the use case. A customer may need the reason an adverse decision occurred; a model-risk team needs reproducibility, validation and monitoring; an operations team needs failure modes and a rollback path. One generic 'AI explanation' cannot serve all three.

A useful rule: as reversibility falls and customer impact rises, human oversight and evidence requirements should rise with it.

09 / PRACTICAL CONTROL CHECKLIST

Before automating a financial decision

Use this page as a research worksheet. The goal is not to slow every deployment; it is to make authority, evidence and recovery explicit before the system goes live.

CONTROL THE WHOLE DECISION LOOP

Do not evaluate only model accuracy. Evaluate the complete decision system: data, model, tools, people, controls, interfaces, monitoring and recovery.

Use the fields below to make task, data, consequence, ownership, explanation and recovery explicit before deployment.

1 / USE CASE + MODEL TASK

2 / DATA SOURCE + FRESHNESS

3 / WRONG-OUTPUT CONSEQUENCE

4 / DECISION OWNER + OVERRIDE

5 / REQUIRED EXPLANATION

6 / MONITORING + ROLLBACK TRIGGER

☐ Bias / subgroup tests☐ Drift / edge-case tests☐ Human review path☐ Customer redress path

A production AI system is a socio-technical control loop. Governance belongs around the whole loop, not only around the model.

Primary sources first

Every material claim in this issue traces to a regulator document or a clearly labelled company annual report.

PRIMARY SOURCE REGISTER

REGULATORS FIRST · COMPANY METRICS LABELLED

S1	Report of the Committee to develop a Framework for Responsible and Ethical Enablement of Artificial Intelligence (FREE-AI) in the Financial Sector 13 August 2025 / RBI committee report. It sets 7 foundational Sutras, 26 recommendations and six strategic pillars for responsible AI adoption in finance.	S6	SEBI (Intermediaries) (Amendment) Regulations, 2025 - responsibility for AI use 6 February 2025 / Final SEBI regulation making covered intermediaries solely responsible for data privacy/security/integrity, AI outputs they rely on, and compliance with law.
S2	Reserve Bank of India Annual Report 2024-25 - Regulation, Supervision and Financial Stability 29 May 2025 / Primary RBI annual-report chapter covering AI in complaint management, supervisory technology, fraud detection and the FREE-AI programme.	S7	SEBI Consultation Paper - Guidelines for Responsible Usage of AI/ML in Indian Securities Markets 20 June 2025 / Primary SEBI consultation paper documenting current market use cases and proposed responsible-AI principles. It is a consultation paper, not final regulation.
S3	RBI Bulletin - Developmental and Regulatory Policies: FREE-AI and MuleHunter.AI December 2024 / Primary RBI policy statement describing the MuleHunter.AI pilot and the rationale for a responsible-AI framework.	S8	SEBI Guidelines for Investment Advisers 8 January 2025 / Primary SEBI circular requiring advisers using AI tools to remain responsible for client data, information used, AI-based advice and legal compliance, and to disclose the extent of AI use.
S4	RBI draft Guidance on Regulatory Principles for Model Risk Management 24 June 2026 / Primary RBI press release for draft guidance. It applies to all models used by covered regulated entities, including third-party and AI/ML models. The guidance was released for public comment and is not presented here as final regulation.	S9	Bajaj Finserv Limited - 19th Annual Report 2025-26 FY2025-26 annual report / Company-reported operating metrics for AI-enabled servicing, lending, document processing, voice-to-data conversion and claims automation. Metrics are reported by Bajaj Finserv and are not independently verified by MarketDeck.
S5	SEBI Board Memorandum - Assigning responsibility for use of AI tools by regulated persons 18 December 2024 / Primary SEBI board material listing observed AI/ML use cases across brokers, depositories, market infrastructure institutions and mutual funds.	S10	Axis Bank Integrated Annual Report 2024-25 FY2024-25 annual report / Company-reported deployment of Axis Deep Intelligence, a GenAI internal assistant across 5,500+ branches supporting 100,000+ employees.

How evidence was handled	RBI and SEBI publications are treated as primary policy sources. Company annual reports are used for company-reported operating metrics and are not independently verified by MarketDeck. No company metric is used as proof that one AI system is superior to another.
Status labels are part of the evidence	A committee report, consultation paper and draft guidance are identified as such. They are not described as final binding rules unless the cited source is a final regulation or circular.
No investment conclusion	This issue evaluates operating use cases, controls and governance. It does not recommend securities, lenders, banks, fintechs or AI vendors and does not estimate investment returns from AI adoption.